

```
In [1]: # Let's import the necessary library.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: # Let's remove the unnecessary warnings.
import warnings
warnings.filterwarnings("ignore")
```

```
In [4]: # Now importing the dataset for the further operation.
customer_details = pd.read_csv("Hospitalisation details.csv")
medical_details = pd.read_csv("Medical Examinations.csv")
customer_name = pd.read_excel("Names.xlsx")
```

```
In [5]: customer_details.shape
```

```
Out[5]: (2343, 9)
```

```
In [6]: medical_details.head()
```

```
Out[6]:
```

	Customer ID	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	NumberOfMajorSurgeries	smoker
0	Id1	47.410	7.47	No	No	No	No major surgery	yes
1	Id2	30.360	5.77	No	No	No	No major surgery	yes
2	Id3	34.485	11.87	yes	No	No	2	yes
3	Id4	38.095	6.05	No	No	No	No major surgery	yes
4	Id5	35.530	5.45	No	No	No	No major surgery	yes

```
In [7]: medical_details.shape
```

```
Out[7]: (2335, 8)
```

```
In [9]: customer_name.head()
```

```
Out[9]:
```

	Customer ID	name
0	Id1	Hawks, Ms. Kelly
1	Id2	Lehner, Mr. Matthew D
2	Id3	Lu, Mr. Phil
3	Id4	Osborne, Ms. Kelsey
4	Id5	Kadala, Ms. Kristyn

```
In [10]: customer_name.shape
```

```
Out[10]: (2335, 2)
```

## Project Task: Week 1

## 1. Collate the files so that all the information is in one place

```
In [11]: # Now combining the data so that all information could be examine in once go through
customer_df1 = pd.merge(customer_name, customer_details, on = "Customer ID")
customer_df1.head()
```

```
Out[11]:
```

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	State ID
0	Id1	Hawks, Ms. Kelly	1968	Oct	12	0	63770.43	tier - 1	tier - 3	R1013
1	Id2	Lehner, Mr. Matthew D	1977	Jun	8	0	62592.87	tier - 2	tier - 3	R1013
2	Id3	Lu, Mr. Phil	1970	?	11	3	60021.40	tier - 1	tier - 1	R1012
3	Id4	Osborne, Ms. Kelsey	1991	Jun	6	1	58571.07	tier - 1	tier - 3	R1024
4	Id5	Kadala, Ms. Kristyn	1989	Jun	19	0	55135.40	tier - 1	tier - 2	R1012

```
In [12]: # Now Lets combine the Last data set and Complete the all information.
final_df = pd.merge(customer_df1, medical_details, on = "Customer ID")
final_df.head()
```

```
Out[12]:
```

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI	HB
0	Id1	Hawks, Ms. Kelly	1968	Oct	12	0	63770.43	tier - 1	tier - 3	R1013	47.410	
1	Id2	Lehner, Mr. Matthew D	1977	Jun	8	0	62592.87	tier - 2	tier - 3	R1013	30.360	
2	Id3	Lu, Mr. Phil	1970	?	11	3	60021.40	tier - 1	tier - 1	R1012	34.485	1
3	Id4	Osborne, Ms. Kelsey	1991	Jun	6	1	58571.07	tier - 1	tier - 3	R1024	38.095	
4	Id5	Kadala, Ms. Kristyn	1989	Jun	19	0	55135.40	tier - 1	tier - 2	R1012	35.530	

```
In [13]: final_df.shape
```

```
Out[13]: (2335, 17)
```

## 2. Check for missing values in the dataset

```
In [14]: final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2335 entries, 0 to 2334
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Customer ID           2335 non-null   object
1   name                   2335 non-null   object
2   year                   2335 non-null   object
3   month                  2335 non-null   object
4   date                   2335 non-null   int64
5   children                2335 non-null   int64
6   charges                 2335 non-null   float64
7   Hospital tier           2335 non-null   object
8   City tier                2335 non-null   object
9   State ID               2335 non-null   object
10  BMI                     2335 non-null   float64
11  HBA1C                   2335 non-null   float64
12  Heart Issues            2335 non-null   object
13  Any Transplants         2335 non-null   object
14  Cancer history          2335 non-null   object
15  NumberOfMajorSurgeries  2335 non-null   object
16  smoker                  2335 non-null   object
dtypes: float64(3), int64(2), object(12)
memory usage: 328.4+ KB
```

```
In [15]: final_df.dtypes.value_counts()
```

```
Out[15]: object      12
float64      3
int64         2
dtype: int64
```

```
In [16]: # Missing values in the data set.
final_df.isnull().sum()
```

```
Out[16]: Customer ID      0
name                  0
year                  0
month                 0
date                  0
children              0
charges               0
Hospital tier         0
City tier              0
State ID              0
BMI                   0
HBA1C                 0
Heart Issues          0
Any Transplants       0
Cancer history        0
NumberOfMajorSurgeries 0
smoker                0
dtype: int64
```

### 3. Find the percentage of rows that have trivial value (for example, ?), and delete such rows if they do not contain significant information

```
In [17]: trivial_value = final_df[final_df.eq("?").any(1)]
trivial_value
```

```
Out[17]:
```

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI
<b>2</b>	Id3	Lu, Mr. Phil	1970	?	11	3	60021.40	tier - 1	tier - 1	R1012	34.485
<b>169</b>	Id170	Torphy, Mr. Bobby	2000	Sep	5	1	37165.16	tier - 1	tier - 3	?	37.620
<b>559</b>	Id560	Pearlman, Mr. Oz	1994	Jul	1	3	17663.14	tier - 1	tier - 3	R1013	23.980
<b>634</b>	Id635	Bruns, Mr. Zachary T	2004	Jul	17	0	15518.18	tier - 2	tier - 3	R1015	25.175
<b>1285</b>	Id1286	Ainsley, Ms. Katie M.	?	Dec	12	1	8547.69	tier - 2	tier - 1	R1013	29.370
<b>1288</b>	Id1289	Levine, Ms. Annie J.	?	Jul	24	0	8534.67	tier - 2	tier - 3	R1024	24.320
<b>1792</b>	Id1793	Capriolo, Mr. Michael	1995	Dec	1	3	4827.90	tier - 1	tier - 2	?	18.905
<b>2317</b>	Id2318	Gagnon, Ms. Candice M	1996	?	18	0	770.38	tier - 3	?	R1012	18.820
<b>2321</b>	Id2322	Street, Ms. Holly	2002	?	19	0	750.00	tier - 3	tier - 1	R1012	21.380
<b>2323</b>	Id2324	Duffy, Ms. Meghan K	1999	Dec	26	0	700.00	?	tier - 3	R1013	22.240

```
In [18]: trivial_value.shape
```

```
Out[18]: (10, 17)
```

```
In [19]: # Percentage of row that have the trivial values
round(trivial_value.shape[0]/final_df.shape[0]*100, 2)
```

```
Out[19]: 0.43
```

```
In [20]: # Now Lets drop the all row that contain the trivial values in the data set.
final_df.drop(final_df[final_df.eq("?").any(1)].index, axis=0, inplace=True)
```

```
In [21]: final_df.shape
```

```
Out[21]: (2325, 17)
```

## 4. Use the necessary transformation methods to deal with the nominal and ordinal categorical variables in the dataset

```
In [22]: # Handling with nominal categorical variable.  
final_df["Heart Issues"].value_counts()
```

```
Out[22]: No      1405  
yes       920  
Name: Heart Issues, dtype: int64
```

```
In [23]: final_df["Any Transplants"].value_counts()
```

```
Out[23]: No      2183  
yes       142  
Name: Any Transplants, dtype: int64
```

```
In [24]: final_df["Cancer history"].value_counts()
```

```
Out[24]: No      1934  
Yes       391  
Name: Cancer history, dtype: int64
```

```
In [25]: final_df["smoker"].value_counts()
```

```
Out[25]: No      1839  
yes       486  
Name: smoker, dtype: int64
```

```
In [26]: # We have some categorical values so first of all we have to transform then by using  
from sklearn.preprocessing import LabelEncoder
```

```
In [27]: le = LabelEncoder()
```

```
In [28]: final_df["Heart Issues"] = le.fit_transform(final_df["Heart Issues"])  
final_df["Any Transplants"] = le.fit_transform(final_df["Any Transplants"])  
final_df["Cancer history"] = le.fit_transform(final_df["Cancer history"])  
final_df["smoker"] = le.fit_transform(final_df["smoker"])
```

```
In [29]: final_df["Heart Issues"].value_counts()
```

```
Out[29]: 0      1405  
1       920  
Name: Heart Issues, dtype: int64
```

```
In [30]: final_df.head()
```

Out[30]:

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI	HB
0	Id1	Hawks, Ms. Kelly	1968	Oct	12	0	63770.43	tier - 1	tier - 3	R1013	47.410	
1	Id2	Lehner, Mr. Matthew D	1977	Jun	8	0	62592.87	tier - 2	tier - 3	R1013	30.360	
3	Id4	Osborne, Ms. Kelsey	1991	Jun	6	1	58571.07	tier - 1	tier - 3	R1024	38.095	
4	Id5	Kadala, Ms. Kristyn	1989	Jun	19	0	55135.40	tier - 1	tier - 2	R1012	35.530	
5	Id6	Baker, Mr. Russell B.	1962	Aug	4	0	52590.83	tier - 1	tier - 3	R1011	32.800	



In [32]:

```
# Handling ordinal categorical variable.  
def clean_ordinal_variable(val):  
    return int(val.replace("tier", "").replace(" ", "").replace("-", ""))
```

In [33]:

```
final_df["Hospital tier"] = final_df["Hospital tier"].map(clean_ordinal_variable)  
final_df["City tier"] = final_df["City tier"].map(clean_ordinal_variable)
```

In [34]:

```
final_df["City tier"].value_counts()
```

Out[34]:

2807

3789

1729

Name: City tier, dtype: int64

In [35]:

```
final_df.head()
```

Out[35]:

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	State ID	BMI	HB
0	Id1	Hawks, Ms. Kelly	1968	Oct	12	0	63770.43	1	3	R1013	47.410	
1	Id2	Lehner, Mr. Matthew D	1977	Jun	8	0	62592.87	2	3	R1013	30.360	
3	Id4	Osborne, Ms. Kelsey	1991	Jun	6	1	58571.07	1	3	R1024	38.095	
4	Id5	Kadala, Ms. Kristyn	1989	Jun	19	0	55135.40	1	2	R1012	35.530	
5	Id6	Baker, Mr. Russell B.	1962	Aug	4	0	52590.83	1	3	R1011	32.800	

# 5. The dataset has State ID, which has around 16 states. All states are not represented in equal proportions in the data. Creating dummy variables for all regions may also result in too many insignificant predictors. Nevertheless, only R1011, R1012, and R1013 are worth investigating further. Create a suitable strategy to create dummy variables with these restraints.

```
In [36]: final_df["State ID"].value_counts()
```

```
Out[36]: R1013      609
          R1011      574
          R1012      572
          R1024      159
          R1026       84
          R1021       70
          R1016       64
          R1025       40
          R1023       38
          R1017       36
          R1019       26
          R1022       14
          R1014       13
          R1015       11
          R1018        9
          R1020        6
          Name: State ID, dtype: int64
```

```
In [37]: Dummies = pd.get_dummies(final_df["State ID"], prefix= "State_ID")
```

```
In [38]: Dummies
```

```
Out[38]:
```

	State_ID_R1011	State_ID_R1012	State_ID_R1013	State_ID_R1014	State_ID_R1015	State_ID_F
0	0	0	1	0	0	
1	0	0	1	0	0	
3	0	0	0	0	0	
4	0	1	0	0	0	
5	1	0	0	0	0	
...	...	...	...	...	...	...
2330	0	0	1	0	0	
2331	0	0	1	0	0	
2332	0	0	1	0	0	
2333	0	0	1	0	0	
2334	0	0	1	0	0	

2325 rows × 16 columns



```
In [39]: # Lets take only those state id which play significant role in the data set.
          Dummy = Dummies[['State_ID_R1011', 'State_ID_R1012', 'State_ID_R1013']]
          Dummy
```



Out[39]:

	State_ID_R1011	State_ID_R1012	State_ID_R1013
0	0	0	1
1	0	0	1
3	0	0	0
4	0	1	0
5	1	0	0
...	...	...	...
2330	0	0	1
2331	0	0	1
2332	0	0	1
2333	0	0	1
2334	0	0	1

2325 rows × 3 columns

In [40]:

```
final_df = pd.concat([final_df, Dummy], axis=1)
```

In [41]:

```
final_df.drop(['State ID'], inplace=True, axis=1)
```

In [42]:

```
final_df.head()
```

Out[42]:

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	BMI	HBA1C	H Is:
0	Id1	Hawks, Ms. Kelly	1968	Oct	12	0	63770.43	1	3	47.410	7.47	
1	Id2	Lehner, Mr. Matthew D	1977	Jun	8	0	62592.87	2	3	30.360	5.77	
3	Id4	Osborne, Ms. Kelsey	1991	Jun	6	1	58571.07	1	3	38.095	6.05	
4	Id5	Kadala, Ms. Kristyn	1989	Jun	19	0	55135.40	1	2	35.530	5.45	
5	Id6	Baker, Mr. Russell B.	1962	Aug	4	0	52590.83	1	3	32.800	6.59	

6. The variable NumberOfMajorSurgeries also appears to have string values. Apply a suitable method to clean up this variable.

```
In [44]: final_df['NumberOfMajorSurgeries'].value_counts()
```

```
Out[44]: No major surgery    1070
1                961
2                272
3                 22
Name: NumberOfMajorSurgeries, dtype: int64
```

```
In [45]: final_df['NumberOfMajorSurgeries'] = final_df['NumberOfMajorSurgeries'].replace('No
```

```
In [46]: final_df['NumberOfMajorSurgeries'] = final_df["NumberOfMajorSurgeries"].astype(int
```

## 7. Age appears to be a significant factor in this analysis. Calculate the patients' ages based on their dates of birth.

```
In [47]: final_df["year"] = pd.to_datetime(final_df["year"], format='%Y').dt.year
final_df["year"]
```

```
Out[47]: 0      1968
1      1977
3      1991
4      1989
5      1962
...
2330    1998
2331    1992
2332    1993
2333    1992
2334    1992
Name: year, Length: 2325, dtype: int64
```

```
In [48]: final_df["month"] = pd.to_datetime(final_df["month"], format='%b').dt.month
final_df["month"]
```

```
Out[48]: 0      10
1       6
3       6
4       6
5       8
..
2330     7
2331     9
2332     6
2333    11
2334     7
Name: month, Length: 2325, dtype: int64
```

```
In [49]: final_df['DateInt'] = final_df["year"].astype(str) + final_df["month"].astype(str)
```


```
In [51]: final_df['DOB'] = pd.to_datetime(final_df.DateInt, format = "%Y%m%d")
```

```
In [52]: final_df.drop(["DateInt"], inplace = True, axis=1)
```

```
In [53]: final_df.head()
```

Out[53]:

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	BMI	HBA1C	H Is:
0	Id1	Hawks, Ms. Kelly	1968	10	12	0	63770.43	1	3	47.410	7.47	
1	Id2	Lehner, Mr. Matthew D	1977	6	8	0	62592.87	2	3	30.360	5.77	
3	Id4	Osborne, Ms. Kelsey	1991	6	6	1	58571.07	1	3	38.095	6.05	
4	Id5	Kadala, Ms. Kristyn	1989	6	19	0	55135.40	1	2	35.530	5.45	
5	Id6	Baker, Mr. Russell B.	1962	8	4	0	52590.83	1	3	32.800	6.59	



In [54]:

```
import datetime as dt
current_date = dt.datetime.now()
```

In [55]:

```
final_df['age'] = (((current_date - final_df.DOB).dt.days)/365).astype(int)
```


In [56]:

```
final_df.head()
```

Out[56]:

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	BMI	...	Heart Issues
0	Id1	Hawks, Ms. Kelly	1968	10	12	0	63770.43	1	3	47.410	...	0
1	Id2	Lehner, Mr. Matthew D	1977	6	8	0	62592.87	2	3	30.360	...	0
3	Id4	Osborne, Ms. Kelsey	1991	6	6	1	58571.07	1	3	38.095	...	0
4	Id5	Kadala, Ms. Kristyn	1989	6	19	0	55135.40	1	2	35.530	...	0
5	Id6	Baker, Mr. Russell B.	1962	8	4	0	52590.83	1	3	32.800	...	0

5 rows × 21 columns



8. The gender of the patient may be an important factor in determining the cost of hospitalization. The salutations in a beneficiary's name can be used to determine their gender. Make a new field for the beneficiary's gender.

```
In [57]: def gender(val):
        if "Ms." in val:
            return 0
        else:
            return 1
```

```
In [58]: final_df["gender"] = final_df["name"].map(gender)
```

```
In [59]: final_df.head()
```


```
Out[59]:
```

	Customer ID	name	year	month	date	children	charges	Hospital tier	City tier	BMI	...	Transpl
0	Id1	Hawks, Ms. Kelly	1968	10	12	0	63770.43	1	3	47.410	...	
1	Id2	Lehner, Mr. Matthew D	1977	6	8	0	62592.87	2	3	30.360	...	
3	Id4	Osborne, Ms. Kelsey	1991	6	6	1	58571.07	1	3	38.095	...	
4	Id5	Kadala, Ms. Kristyn	1989	6	19	0	55135.40	1	2	35.530	...	
5	Id6	Baker, Mr. Russell B.	1962	8	4	0	52590.83	1	3	32.800	...	


5 rows × 22 columns

9. You should also visualize the distribution of costs using a histogram, box and whisker plot, and swarm plot.


```
In [61]: # Histogram for the cost distribution.
plt.figure(figsize=(12,5))
sns.histplot(final_df['charges'])
plt.show()
```



```
In [62]: # Visualize the cost distribution of the hospitals by box or whisker plot.
plt.boxplot(final_df['charges'])
plt.show()
```




```
In [63]: # Visualize the cost distribution of the hospitals by swarm plot.
plt.figure(figsize=(12,5))
sns.swarmplot(x='year', y='charges', hue="gender", data=final_df)
plt.xticks(rotation=90)
plt.show()
```



## 10. State how the distribution is different across gender and tiers of hospitals

```
In [64]: sns.countplot(data = final_df, x='Hospital tier', hue= 'gender')
plt.show()
```



## 11. Create a radar chart to showcase the median hospitalization cost for each tier of hospitals

```
In [65]: print("median cost of tier 1 hospitals:", final_df[final_df["Hospital tier"]==1].cost.median())
print("median cost of tier 2 hospitals:", final_df[final_df["Hospital tier"]==2].cost.median())
print("median cost of tier 3 hospitals:", final_df[final_df["Hospital tier"]==3].cost.median())
```

```
median cost of tier 1 hospitals: 32097.434999999998
median cost of tier 2 hospitals: 7168.76
median cost of tier 3 hospitals: 10676.83
```

```
In [66]: df = pd.DataFrame(dict(r=[32097.43, 7168.76, 10676.83],theta=['tier 1 hospital','tier 2 hospital','tier 3 hospital']))
```

```
In [67]: df
```

```
Out[67]:
```

	r	theta
0	32097.43	tier 1 hospital
1	7168.76	tier 2 hospital
2	10676.83	tier 3 hospital

```
In [68]: import plotly.express as px
fig = px.line_polar(df, r='r', theta='theta', line_close=True)
fig.update_traces(fill='toself')
fig.show()
```



## 12. Create a frequency table and a stacked bar chart to visualize the count of people in the different tiers of cities and hospitals

```
In [69]: # Frequency table for count of the people according to the tier of city and hospital
final_df["Hospital tier"].value_counts()
```

```
Out[69]: 2    1334
         3     691
         1     300
         Name: Hospital tier, dtype: int64
```

```
In [70]: city_freq = final_df["City tier"].value_counts().rename_axis('City&hospital_tier')
```

```
In [71]: hospital_freq = final_df["Hospital tier"].value_counts().rename_axis('City&hospital_tier')
```


```
In [72]: df = pd.merge(city_freq, hospital_freq, on = 'City&hospital_tier')
```

```
In [73]: df
```

```
Out[73]:
```

	City&hospital_tier	city_counts	hospital_counts
0	2	807	1334
1	3	789	691
2	1	729	300

```
In [74]: plt.bar(df["City&hospital_tier"], df["city_counts"], color='r')
plt.bar(df["City&hospital_tier"], df["hospital_counts"], bottom=df["city_counts"],
plt.show()
```



```
In [75]: from scipy.stats import ttest_1samp
```

H0: the distributions of all samples are equal. || H1: the distributions of one or more samples are not equal

```
In [76]: from scipy.stats import friedmanchisquare
data1 = [32097.43]
data2 = [7168.76]
data3 = [10676.83]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=2.000, p=0.368  
Probably the same distribution

```
In [77]: # b. The average hospitalization costs for the three types of cities are not signifi
print("median cost of tier 1 city:", final_df[final_df["City tier"]==1].charges.mean())
print("median cost of tier 2 city:", final_df[final_df["City tier"]==2].charges.mean())
print("median cost of tier 3 city:", final_df[final_df["City tier"]==3].charges.mean())
```

median cost of tier 1 city: 10027.15  
median cost of tier 2 city: 8968.33  
median cost of tier 3 city: 9880.07

```
In [78]: data1 = [10027.15]
data2 = [8968.33]
data3 = [9880.07]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
```



```
print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=2.000, p=0.368  
Probably the same distribution

In [79]: *# c. The average hospitalization cost for smokers is not significantly different from non-smokers*

```
print("median cost of smoker:", final_df[final_df["smoker"]==1].charges.median())
print("median cost of non smoker:", final_df[final_df["smoker"]==0].charges.median())
```

median cost of smoker: 34125.475  
median cost of non smoker: 7537.16

In [80]:

```
from scipy.stats import kruskal
data1 = [34125.475]
data2 = [7537.16]
stat, p = kruskal(data1, data2)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=1.000, p=0.317  
Probably the same distribution

In [81]: *# d. Smoking and heart issues are independent*

```
from scipy.stats import chi2_contingency
table = [[final_df["Heart Issues"].value_counts(), final_df["smoker"].value_counts()],
          [final_df["Heart Issues"].value_counts(), final_df["smoker"].value_counts()]]
stat, p, dof, expected = chi2_contingency(table)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably independent')
else:
    print('Probably dependent')
```

stat=191.145, p=0.000  
Probably dependent

## Project Task: Week 2

1. Examine the correlation between predictors to identify highly correlated predictors. Use a heatmap to visualize this.

In [82]: `final_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2325 entries, 0 to 2334
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          2325 non-null   object
1   name                                2325 non-null   object
2   year                                2325 non-null   int64
3   month                              2325 non-null   int64
4   date                                2325 non-null   int64
5   children                            2325 non-null   int64
6   charges                             2325 non-null   float64
7   Hospital tier                        2325 non-null   int64
8   City tier                            2325 non-null   int64
9   BMI                                 2325 non-null   float64
10  HBA1C                               2325 non-null   float64
11  Heart Issues                         2325 non-null   int32
12  Any Transplants                     2325 non-null   int32
13  Cancer history                      2325 non-null   int32
14  NumberOfMajorSurgeries              2325 non-null   int32
15  smoker                              2325 non-null   int32
16  State_ID_R1011                      2325 non-null   uint8
17  State_ID_R1012                      2325 non-null   uint8
18  State_ID_R1013                      2325 non-null   uint8
19  DOB                                 2325 non-null   datetime64[ns]
20  age                                 2325 non-null   int32
21  gender                              2325 non-null   int64
dtypes: datetime64[ns](1), float64(3), int32(6), int64(7), object(2), uint8(3)
memory usage: 315.6+ KB
```

```
In [83]: # In the data frame same of the column are not usable to model building so Lets fi
#then indentify the highly corelated predictor.
final_df.drop(["Customer ID", 'name', 'year', 'month', 'date', 'DOB'], inplace=True)
final_df.shape
```

Out[83]: (2325, 16)

```
In [84]: final_df.head()
```


Out[84]:

	children	charges	Hospital tier	City tier	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	NumberOfMajo
0	0	63770.43	1	3	47.410	7.47	0	0	0	
1	0	62592.87	2	3	30.360	5.77	0	0	0	
3	1	58571.07	1	3	38.095	6.05	0	0	0	
4	0	55135.40	1	2	35.530	5.45	0	0	0	
5	0	52590.83	1	3	32.800	6.59	0	0	0	


```
In [85]: corr = final_df.corr()
corr
```

Out[85]:

	children	charges	Hospital tier	City tier	BMI	HBA1C	Heart Issues
children	1.000000	0.055901	-0.052438	-0.015760	-0.005339	-0.101379	0.023984
charges	0.055901	1.000000	-0.446687	0.035300	0.346730	0.139697	0.049299
Hospital tier	-0.052438	-0.446687	1.000000	-0.039755	-0.104771	0.057855	0.053376
City tier	-0.015760	0.035300	-0.039755	1.000000	0.038123	-0.005404	0.023152
BMI	-0.005339	0.346730	-0.104771	0.038123	1.000000	-0.006920	0.017129
HBA1C	-0.101379	0.139697	0.057855	-0.005404	-0.006920	1.000000	0.007699
Heart Issues	0.023984	0.049299	0.053376	0.023152	0.017129	0.007699	1.000000
Any Transplants	-0.142040	-0.127028	0.011729	0.002970	0.015893	-0.159855	-0.140260
Cancer history	-0.027880	-0.022522	-0.021429	-0.018639	-0.020235	-0.170921	0.111195
NumberOfMajorSurgeries	-0.113161	0.053308	0.033230	0.027937	0.018851	-0.091594	0.206140
smoker	0.017713	0.838462	-0.474077	0.032034	0.107126	0.007257	-0.007150
State_ID_R1011	0.011666	0.286956	-0.114685	0.036049	0.115671	0.015525	0.005850
State_ID_R1012	0.005247	-0.074636	0.020272	-0.018253	0.017939	-0.019513	0.021770
State_ID_R1013	-0.013834	-0.150634	0.002455	0.002766	-0.208744	0.033453	-0.027960
age	-0.005457	0.304395	0.133771	-0.008070	0.049260	0.460558	0.192270
gender	0.011205	0.034069	0.041261	0.054073	0.079930	-0.027339	0.010270



```
In [86]: plt.figure(figsize=(15,10))
sns.heatmap(corr, annot=True, linewidth=.5, cmap="crest")
plt.show()
```



## 2. Develop and evaluate the final model using regression with a stochastic gradient descent optimizer. Also, ensure that you apply all the following suggestions:

Note: • Perform the stratified 5-fold cross-validation technique for model building and validation • Use standardization and hyperparameter tuning effectively • Use sklearn-pipelines • Use appropriate regularization techniques to address the bias-variance trade-off

- Create five folds in the data, and introduce a variable to identify the folds
- For each fold, run a for loop and ensure that 80 percent of the data is used to train the model and the remaining 20 percent is used to validate it in each iteration
- Develop five distinct models and five distinct validation scores (root mean squared error values)
- Determine the variable importance scores, and identify the redundant variables

```
In [87]: # Lets first separate the input and output data.
x = final_df.drop(["charges"], axis=1)
y = final_df[["charges"]]
```

```
In [88]: # Lets split the data set into the training and testing data.
from sklearn.model_selection import train_test_split
```

```
In [89]: x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=.20, random_state=42)
```

```

In [90]: # Now standardize the data.
         from sklearn.preprocessing import StandardScaler

In [91]: sc = StandardScaler()

In [92]: x_train = sc.fit_transform(x_train)
         x_test = sc.fit_transform(x_test)

In [93]: from sklearn.linear_model import SGDRegressor

In [94]: from sklearn.model_selection import GridSearchCV

         params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5,
                             0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0,
                             9.0, 10.0, 20, 50, 100, 500, 1000],
                    'penalty': ['l2', 'l1', 'elasticnet']}

         sgd = SGDRegressor()

         # Cross Validation
         folds = 5
         model_cv = GridSearchCV(estimator = sgd,
                                  param_grid = params,
                                  scoring = 'neg_mean_absolute_error',
                                  cv = folds,
                                  return_train_score = True,
                                  verbose = 1)
         model_cv.fit(x_train, y_train)

         Fitting 5 folds for each of 84 candidates, totalling 420 fits
Out[94]: GridSearchCV(cv=5, estimator=SGDRegressor(),
                    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
                                           0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                                           4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
                                           100, 500, 1000],
                                'penalty': ['l2', 'l1', 'elasticnet']}},
                    return_train_score=True, scoring='neg_mean_absolute_error',
                    verbose=1)

In [95]: model_cv.best_params_

Out[95]: {'alpha': 50, 'penalty': 'l1'}

In [96]: sgd = SGDRegressor(alpha=100, penalty='l1')

In [97]: sgd.fit(x_train, y_train)

Out[97]: SGDRegressor(alpha=100, penalty='l1')

In [98]: sgd.score(x_test, y_test)

Out[98]: 0.8602495677726669

In [99]: y_pred = sgd.predict(x_test)

In [100... from sklearn.metrics import mean_squared_error, mean_absolute_error

In [102... sgd_mae = mean_absolute_error(y_test, y_pred)
          sgd_mse = mean_squared_error(y_test, y_pred)

```

```
sgd_rmse = sgd_mse*(1/2.0)
```

```
In [103... print("MAE:", sgd_mae)
print("MSE:", sgd_mse)
print("RMSE:", sgd_rmse)
```

```
MAE: 3134.524710884731
MSE: 23506849.75240736
RMSE: 11753424.87620368
```

```
In [104... # d. Determine the variable importance scores, and identify the redundant variables
importance = sgd.coef_
```

```
In [105... pd.DataFrame(importance, index = x.columns, columns=['Feature_imp'])
```

```
Out[105]:
```

	Feature_imp
children	342.564040
Hospital tier	-1179.487983
City tier	0.000000
BMI	2718.546547
HBA1C	60.773832
Heart Issues	0.000000
Any Transplants	0.000000
Cancer history	19.557174
NumberOfMajorSurgeries	0.000000
smoker	8793.400707
State_ID_R1011	-189.577264
State_ID_R1012	0.000000
State_ID_R1013	-350.073301
age	3353.637397
gender	0.000000

### 3. Use random forest and extreme gradient boosting for cost prediction, share your crossvalidation results, and calculate the variable importance scores

random forest

```
In [106... from sklearn.ensemble import RandomForestRegressor
```

```
In [107... # Instantiate model with 1000 decision trees
rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)

# Train the model on training data
rf.fit(x_train, y_train)
```

```
Out[107]: RandomForestRegressor(n_estimators=1000, random_state=42)
```

```
In [108... score = rf.score(x_test,y_test)
score
```

```
Out[108]: 0.9222696338245824
```

```
In [109... y_pred = rf.predict(x_test)
```

```
In [110... rf_mae = mean_absolute_error(y_test, y_pred)
```

```
In [111... from sklearn.ensemble import GradientBoostingRegressor
```

```
In [112... # Instantiate model with 1000 decision trees
gbr = GradientBoostingRegressor(n_estimators = 1000, random_state = 42)

# Train the model on training data
gbr.fit(x_train, y_train)
```

```
Out[112]: GradientBoostingRegressor(n_estimators=1000, random_state=42)
```

```
In [113... score = gbr.score(x_test,y_test)
score
```

```
Out[113]: 0.9042734212625119
```

```
In [114... y_pred = gbr.predict(x_test)
```

```
In [115... gbr_mae = mean_absolute_error(y_test, y_pred)
gbr_mae
```

```
Out[115]: 2375.8700944163274
```

## 4. Case scenario:

Estimate the cost of hospitalization for Christopher, Ms. Jayna (her date of birth is 12/28/1988, height is 170 cm, and weight is 85 kgs). She lives in a tier-1 city and her state's State ID is R1011. She lives with her partner and two children. She was found to be nondiabetic (HbA1c = 5.8). She smokes but is otherwise healthy. She has had no transplants or major surgeries. Her father died of lung cancer. Hospitalization costs will be estimated using tier-1 hospitals.

```
In [116... # Calculate the age of the person.
date = str(19881228)
date1 = pd.to_datetime(date, format = "%Y%m%d")
```

```
In [117... current_date = dt.datetime.now()
current_date
```

```
Out[117]: datetime.datetime(2023, 5, 7, 23, 29, 41, 979602)
```

```
In [119... age = (current_date - date)
age
```

```
-----
TypeError                                Traceback (most recent call last)
Input In [119], in <cell line: 1>()
----> 1 age = (current_date - date)
      2 age

TypeError: unsupported operand type(s) for -: 'datetime.datetime' and 'str'
```

```
In [120... age = int(12421/365)
age
```

```
Out[120]: 34
```

```
In [121... # now with the help of height and weight we will calculate the BMI.
height_m = 170/100
height_sq = height_m*height_m
BMI = 85/height_sq
np.round(BMI,2)
```

```
Out[121]: 29.41
```

```
In [122... # Now Lets gen
list = [[2,1,1,24.41,5.8,0,0,0,0,1,1,0,0,34,0]]
```

```
In [123... df = pd.DataFrame(list, columns = ['children', 'Hospital tier', 'City tier', 'BMI',
                                     'Cancer history', 'NumberOfMajorSurgeries', 'smoker',
                                     'State_ID_R1013', 'age', 'gender'] )
df
```

```
Out[123]:
```

	children	Hospital tier	City tier	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	NumberOfMajorSurgeries
0	2	1	1	24.41	5.8	0	0	0	0

## 5. Find the predicted hospitalization cost using all models. The predicted value should be the mean of the five models' predicted values.

```
In [124... Hospital_cost = []
# Now Lets predict the hospitalization cost through SGDRegressor
Cost1 = sgd.predict(df)
Hospital_cost.append(Cost1)
# Now Lets predict the hospitalization cost through Random Forest
Cost2 = rf.predict(df)
Hospital_cost.append(Cost2)
# Now Lets predict the hospitalization cost through Extreme gradient Booster
Cost3 = gbr.predict(df)
Hospital_cost.append(Cost3)
avg_cost = np.mean(Hospital_cost)
avg_cost
```

```
Out[124]: 103999.74746407126
```

So in the new case the avg predicted hospitalization cost is 104922.59



In [ ]: